

# GPU Programming



Farshad Khunjush

Department of Computer Science and Engineering  
Shiraz University  
Fall 2025

# Introduction to GPUs

## Part 1




Farshad Khunjush


Slides come from  
Professor Babak Falsafi @ EPFL

# Multithreading (Review)

- Pipelines waste cycles in two ways
- **Vertical**: whole cycle empty, nothing issued
  - Most common after long latency events
- **Horizontal**: unable to use full issue width
  - Software not exposing enough ILP for hardware




# Coarse Grain Multithreading




- Switch to a new thread on **a long latency event**
  - Pay a small cost to switch in a new context
  - Addresses purely vertical waste

# CGMT Performance




- **Critical decision:** when to switch threads
  - When current thread's utilization is about to drop
  - (e.g. L2 cache miss)
- **Requirements for improving throughput:**
  - (Thread switch) + (pipe fill time) << blocking latency
    - Need useful work to be done before other thread comes back
  - Fast thread-switch: **multiple register banks**
  - Fast pipe-fill: **short pipe**
- **Advantage:** small changes to existing hardware
- **Drawback:** single-thread performance suffers

# Fine Grained Multithreading (FGMT)




- ☐ Cycle between multiple threads periodically
  - Eliminate “switch time” by keeping all thread state hot

# FGMT vs. CGMT



- **Thread switch policy:**
  - Most historical designs of **FGMT** are round robin
  - **CGMT** swaps on long latency events
- **FGMT** addresses much shorter latencies
  - Hardware costs to keep all contexts immediately ready
- For apps. with abundant ILP, **FGMT could suffer**
  - Can introduce “flexible interleaving” to solve
  - One thread remains scheduled for many cycles


# Simultaneous Multithreading (SMT)



- Instructions from multiple threads in **same cycle**
  - First design that can **reduce horizontal waste**

# Taking Parallelism Further

- We started with a single instruction CPU...
  - Added contexts, different ways to multithread
  - Makes our processor more latency-tolerant
  - Also added ALUs to make it SIMD capable



# Limitation on Arithmetic Units

---

- Adding ALUs has an upper limit
  - We need enough instructions to utilize all of them
  - Recall last time, pipeline width has diminishing benefit
  
- SMT addresses horizontal waste, by bringing in more independent instructions
  - Also has an upper limit: scaling register files and ROB!

# Better Idea: One Instruction, All ALUs

---

- Previously, rely on programs to expose ILP
  - SMT is a pure hardware technique to exploit it
- Now, control all of the ALUs with one instruction
  - This is a “**vector processor**”, built ages ago  
(Cray 1 - 1976, Illinois ILLIAC IV – 1972)
- Key point: we’ve gone from “program threads” to “**SIMD threads**” running in **lockstep**

# Vector vs. SIMD

---


- SIMD instructions have **configurable width**, and actually move data there
  - e.g., x86 operations call them **xmms**
    - Can specify 64, 128, 256, 512 bit widths
- Vector processors operate in **strides**
  - Gather elements from memory, operate in parallel on all of them
    - No need to change code for different SIMD implementations!
- GPUs are more like vector processors
  - Since instructions are executed all together, operations happen in large vectors

# SIMD Processing

---

- Single instruction operates on multiple data elements
  - In time or in space
- Multiple processing elements (PEs), i.e., execution units
- Time-space duality
  - **Array processor**: Instruction operates on multiple data elements at the **same time** using **different spaces (PEs)**
  - **Vector processor**: Instruction operates on multiple data elements in **consecutive time steps** using the **same space (PE)**


# Array vs. Vector Processors



# Example: Image Processing


---

- Simple operation to do a fade out
  - Take every pixel and multiply by a constant (e.g. 0.7)
- Single thread programming: loop over all pixels




# Example: Image Processing

- Control flow happens **every loop iteration**
  - One instruction at a time
  - **Optimizations** are all done **in the hardware**
    - e.g. Loop unrolling, branch prediction, SMT



# Example: Image Processing

- ❑ But, given hypothetical software parallel language
  - Divide the work among a bunch of threads
  - Do everything in parallel




```
// hypothetically in parallel
for_all_pixels {
    array[i] *= fade;
}
```

# Massively Parallel Hardware


---

- **Scale** pipeline width, e.g., 1k ALUs
- **Key insight:** use a single instruction stream
  - Because we know all threads are doing the same thing  
(e.g. a multiplication operation and memory writes)



# Massively Parallel Hardware

- Works great for **embarrassingly parallel programs**
- **The problem?** Memory and long latency ops.
  - Solution: Use huge degree of multithreading
    - Every clock cycle should have a thread ready to execute




# ~~Massively Parallel Hardware~~ GPUs

---

- Make **cores** simple
  - In-order pipelines
- Put **thousands of ALUs** on the die
  - NVIDIA Volta has ~5.5k integer and ~2.5k FP ALUs
  - Reminder: NVIDIA calls these “**cores**”, but they are not similar to CPU cores
  - Parallel software **favors throughput**
- Share cores among **thousands of threads**
- Take **throughput-latency** trade-off to extreme
  - Trillions of integer operations per second
  - ... but, huge single thread latency

# Example: NVIDIA Tesla V100



**5120 CUDA cores!**  
**(much smaller than CPU cores )**

# Example: NVIDIA Volta V100




22

# Example: CPU vs. GPU

Intel Skylake-X:

- ~600 mm<sup>2</sup>
- 28 cores/56 threads
- Hierarchy:  
32KB/1MB/38.5MB/1.  
5TB
- 205 Watts



GPU V100:

- 815 mm<sup>2</sup>
- 5120 CUDA cores
- 2MB reg file
- Hierarchy:  
96KB/256KB/16GB
- 300 Watts




# One Slide of GPU History

---


- It doesn't seem like it, but the popularity of modern GPUs sort of happened by accident
  - They were originally meant for graphics **only**
  - **No** general purpose programming language
    - Have fun with OpenGL
- Researchers started using them for computation around 2003 – called “General Purpose GPUs” (**GPGPUs**)
  - Plenty of data parallel computation (e.g., simulation)
  - The latest killer app is AI/ML with DNNs
  - Now, all GPUs can run regular declarative programs

# Increase in GPU Popularity

- Remember this graph?
  - Free lunch is over
- Moore's Law ver. 2
  - Chip density still increases
  - Clock speed does not
- GPUs turn density to compute (parallelism)
- CPUs on single thread performance
  - But, this has saturated



# GPU Parallelism Keeps Scaling



# Why are GPUs Interesting?

---

- **Good platform** for today's CS problems
  - If you are interested **in doing ML**, you will use GPUs!
  - **Geoff Hinton** - co-inventor of back propagation algorithm used for training multi-level neural networks
    - ▣ Quoted as saying GPUs are perfect match for deep belief networks (startup bought by Google in 2013)
  - **Applications**: self-driving cars, augmented reality...
- **Great example of HW/SW co-design**
  - **New software** to express parallelism
  - **Hardware designers** can put resources to best use


# When To Use GPUs?

---

- Dependent on your program...
  - They aren't magic!
- Remember fundamental trade-off
  - Multithreading **always** reduces single-thread perf.
  - GPUs take that to the extreme
- GPUs run **regularly** parallel programs well
  - Very common for physics simulation, machine learning
  - Irregular parallelism with synchronization more difficult
    - System is not built for that, but improving every year

# Overview of GPU Architecture

- Essentially high-performance **coprocessors**
- Attached to the system in a few different ways
  - Integrated onto chip with **unified memory**
    - Cache-coherent with CPU (e.g., AMD)
  - **Discrete device**, connected through I/O
    - PCI Express (PCIe) bus



Integrated

We use these!




Discrete

# GPU Architecture Hierarchy

- ❑ Note: we use **NVIDIA** terminology in this class


(because those are the GPUs you will use)

- ❑ Modern GPU (**V100**)
  - 6 processing clusters
  - 14 **SMs** per cluster
  - 32 thread blocks per **SM**
  - 32 threads per **warp**




# Streaming Multiprocessor (SM)

- NVIDIA's equivalent of a processor
  - Broken down into:
    - Front-end (i-cache and scheduler)
    - Registers
    - Execution units/ALUs
    - Memory
- Wide execution path, all threads execute on all cores



# Parallelism Model - SIMT

- **SIMT** = Single Instruction Multiple Thread
  - Remember vector processors
  - **GPU threads** act like the ALUs inside each vector slot
- **SIMT** is the combination of SIMD operations, with scheduling in round robin fashion



# Definition of a Thread

---

- What's a thread for you?
  - On a CPU, a thread has its own execution context
    - Instruction stream, stack, program counter
- On a GPU, **this is not true**
  - All threads **share an execution context** (including stack)
  - All threads in a program have **the same instruction stream**

# Constructing Warps

---

- A **warp** is simple: 32 threads running together
  - Natural choice because the **pipeline** is **32-wide SIMT**
- Choice made to simplify hardware scheduling
  - Otherwise, SMs track thousands of threads
    - Current SMs allow 48 warps per SM, so 1536 CUDA threads
  - Remember what we learned about SMT
    - Tracking memory/register readiness for 1k threads is hard


# Important Corollary of Using Warps

---

- Since **all threads** of a **warp execute together**
  - They all have to be ready, or can't execute!!!
- If one thread out of 32 is waiting... (e.g., memory)
  - SM hardware will not schedule the warp
  - **Huge vertical pipeline waste**
- **H/W solves it with more warps** to hide latency
  - But, better S/W solution, cover next lectures...
  - Also, talk about thread divergence (similar problem)

# Multithreading in GPUs


- GPUs use FGMT to hide long latency
  - Last lectures, the “thread” was a large CPU thread
  - This time, the “threads” are simply different warps



**Warp: group of threads running in lockstep**

# Multithreading in GPUs

- GPUs use FGMT to hide long latency
  - Last lecture, the “thread” was a large CPU thread
  - This time, the “threads” are simply different warps




# Allocating Warps to SMs

---

- Recall:
  - Modern GPU has ~80 SMs, max 48 warps each
- Two levels of resource allocation
- Global allocation maps groups of warps to cores
  - “Groups of warps” = Thread Blocks (stay tuned!)
  - Allocates registers in cores
- Instruction scheduler chooses warps to issue
  - Find warps with independent instructions

# Why GPUs are so Fast: GPUs vs. CPUs



# Fermi vs. Kepler GPUs

	FERMI GF100	FERMI GF104	KEPLER GK104	KEPLER GK110
<b>Compute Capability</b>	2.0	2.1	3.0	3.5
<b>Threads / Warp</b>	32	32	32	32
<b>Max Warps / Multiprocessor</b>	48	48	64	64
<b>Max Threads / Multiprocessor</b>	1536	1536	2048	2048
<b>Max Thread Blocks / Multiprocessor</b>	8	8	16	16
<b>32-bit Registers / Multiprocessor</b>	32768	32768	65536	65536
<b>Max Registers / Thread</b>	63	63	63	255
<b>Max Threads / Thread Block</b>	1024	1024	1024	1024
<b>Shared Memory Size Configurations (bytes)</b>	16K	16K	16K	16K
	48K	48K	32K	32K
<b>Max X Grid Dimension</b>	$2^{16-1}$	$2^{16-1}$	$2^{32-1}$	$2^{32-1}$
<b>Hyper-Q</b>	No	No	No	Yes
<b>Dynamic Parallelism</b>	No	No	No	Yes

Courtesy NVIDIA